



Keypoints and Features

Federico Tombari

CGLibs, Pisa

June 4, 2013



A feature..what?

Etymology

From [Anglo-Norman *feture*](#), from Old French *faiture*, from Latin *factura*.

Pronunciation

- (UK) IPA: /fi:tʃə(ɹ)/, X-SAMPA: /fi:tS@(r)/

- Audio (US)




- Rhymes: -i:tʃə(ɹ)

Noun

feature (*plural* **features**)


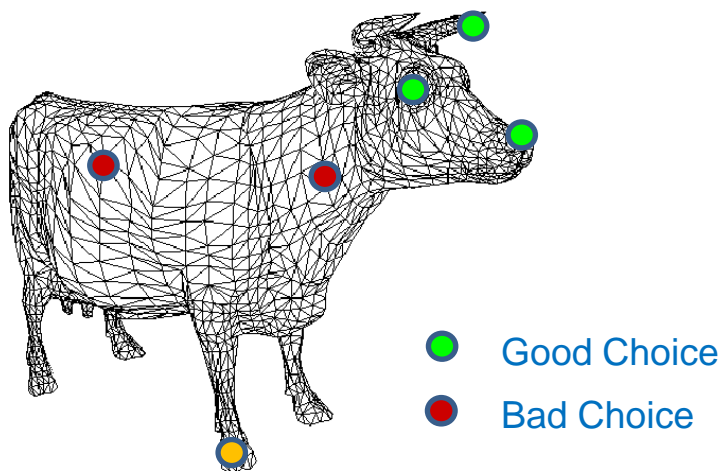
1. (*obsolete*) One's [structure](#) or [make-up](#); [form](#), [shape](#), [bodily proportions](#). [\[quotations ▼\]](#)
2. An important or [main item](#).
3. (*media*) A long, prominent, [article](#) or [item](#) in the [media](#), or the [department](#) that creates them; frequently used technically to [distinguish](#) content from news.
4. Any of the physical constituents of the face ([eyes](#), [nose](#), etc.).
5. (*computing*) A [beneficial capability](#) of a piece of [software](#). [\[quotations ▼\]](#)
6. The cast or structure of anything, or of any part of a thing, as of a [landscape](#), a [picture](#), a [treaty](#), or an [essay](#); any marked peculiarity or characteristic; as, one of the [features](#) of the [landscape](#). [\[quotations ▼\]](#)
7. (*archaeology*) Something discerned from physical evidence that helps define, identify, characterize, and interpret an archeological site. [\[quotations ▼\]](#)
8. (*engineering*) Characteristic [forms](#) or [shapes](#) of a [part](#). For example, a [hole](#), [boss](#), [slot](#), [cut](#), [chamfer](#), or [fillet](#).

- Feature is a compact – but *rich* – *representation* of our (3D) data
- It is designed to be invariant (or robust) to a specific class of transformations and/or set of disturbances






- **3D keypoints** are
 - **Distinctive**, i.e. suitable for effective description and matching (*globally definable*)
 - **Repeatable** with respect to point-of-view variations, noise, etc... (*locally definable*)
- Usually scale-invariance is not an issue (but better if each feature is extracted together with its characteristic scale)



Distinctiveness vs. repeatability

- (for now) a small set of detectors specifically proposed for 3D point clouds and range maps
 - Intrinsic Shape Signatures (ISS) [Zhong ICCVW09]
 - NARF [Steder ICRA11]
 - Uniform Sampling (basically a *voxelGrid*, where selected points are a subset of the input cloud)
- Several detectors «derived» from 2D interest point detectors
 - Harris (2D, 3D, 6D) [Harris AVC88] - CD
 - SIFT [Lowe IJCV04] - BD
 - SUSAN [Smith IJCV95] - CD
 - AGAST [Mair ECCV10] - CD



Courtesy of Unnikrishnan & Hebert

- In the context of most PCL applications scale is not an issue
BUT
- The characteristic scale is still an important property of a 3D keypoint
- Several recent proposals, two main categories [Tombari IJCV13]
 - **Fixed-scale detectors:** all keypoints are detected at a specific scale (input parameter)
 - **Adaptive-scale detectors:** specific scale-space analysis to detect salient structures at multiple scales, associating each keypoint a **characteristic scale**
- Need for performance assessment
 - Locality repeatability / Quantity, Scale repeatability, Efficiency
 - www.vision.deis.unibo.it/keypoints3d

- Exploits the covariance matrix
$$\mathbf{M}(\mathbf{p}_i) = \frac{1}{\sum_{j=1}^k \rho_i} \sum_{j=1}^k \rho_i (\mathbf{p}_j - \mathbf{p}_i) (\mathbf{p}_j - \mathbf{p}_i)^T$$


- Let its eigenvalues, in decreasing magnitude order, be



$$\lambda_1, \lambda_2, \lambda_3$$

- The pruning step discards points with similar spreads along the principal directions, where a repeatable LRF cannot be defined

$$\frac{\lambda_2(\mathbf{p})}{\lambda_1(\mathbf{p})} < Th_{12} \wedge \frac{\lambda_3(\mathbf{p})}{\lambda_2(\mathbf{p})} < Th_{23}$$

- Saliency is the magnitude of the third eigenvalue $\rho(\mathbf{p}) \doteq \lambda_3(\mathbf{p})$
- Non-Maxima Suppression (NMS) over saliency
- It includes only points with large variations along each principal direction
- “Winner” of PCL 3D detector evaluation in [Filipe 2013]










UNIFORM SAMPLING

```
pcl::PointCloud<int> indices;  
pcl::UniformSampling<pcl::PointXYZ> uniform_sampling;  
uniform_sampling.setInputCloud (cloud);  
uniform_sampling.setRadiusSearch (0.05f); //the 3D grid leaf size  
uniform_sampling.compute (indices);
```

ISS

```
pcl::PointCloud<pcl::PointXYZ>::Ptr keypoints (new  
pcl::PointCloud<pcl::PointXYZ>());  
pcl::ISSKeypoint3D<pcl::PointXYZ, pcl::PointXYZ> iss_detector;  
iss_detector.setSalientRadius (support_radius);  
iss_detector.setNonMaxRadius (nms_radius);  
iss_detector.setInputCloud (cloud);  
iss_detector.compute (*keypoints);
```


- 3 orthogonal unit vectors defined upon a local support
- Goal:
 - invariant to rotations and translations
 - robust to noise and clutter
- Common approach to deal with ambiguities in the LRF definition
 - Define **multiple LRFs** at each keypoint, providing multiple descriptions of the same keypoint
 - Cons:
 - more descriptors to be computed and matched (less efficient)
 - ambiguity pushed to the matching stage
 - Eg. EVD of the scatter matrix computed over the support as used in [Mian10] [Novatnack08] [Zhong09], provides 3 repeatable directions but **no repeatable sign** [Tombari10]
 - 4 different RFs can be obtained by enforcing the right-hand rule



```
pcl::PointCloud< pcl::ReferenceFrame >::Ptr lrfs(new
    pcl::PointCloud< pcl::ReferenceFrame> ());

pcl::BOARDLocalReferenceFrameEstimation<pcl::PointXYZ, pcl::Normal,
pcl::ReferenceFrame> lrf_est;

lrf_est.setRadiusSearch (0.5f);
lrf_est.setInputCloud (keypoints);
lrf_est.setInputNormals (cloud_normals);
lrf_est.setSearchSurface (cloud);

lrf_est.compute (*lrfs);
```



- **compact** representations aimed at detecting similarities between surfaces (*surface matching*)
- based on the support size
 - **Pointwise descriptors**
 - Simple, efficient, but not robust to noise, often not descriptive enough
 - **Local/Regional descriptors**
 - Well suited to handle clutter and occlusions
 - Can be vector quantized in codebooks
 - Segmentation, registration, recognition in clutter, 3D SLAM
 - **Global descriptors**
 - Complete information concerning the surface is needed (no occlusions and clutter, unless pre-processing)
 - Higher invariance, well suited for **retrieval and categorization**
 - More descriptive on objects with poor geometric structure (household objects..)





- Spin Image descriptor [Johnson99] is arguably the most popular 3D local descriptor
- 2D histograms accumulating points by spinning around a repeatable axis (*normal*)



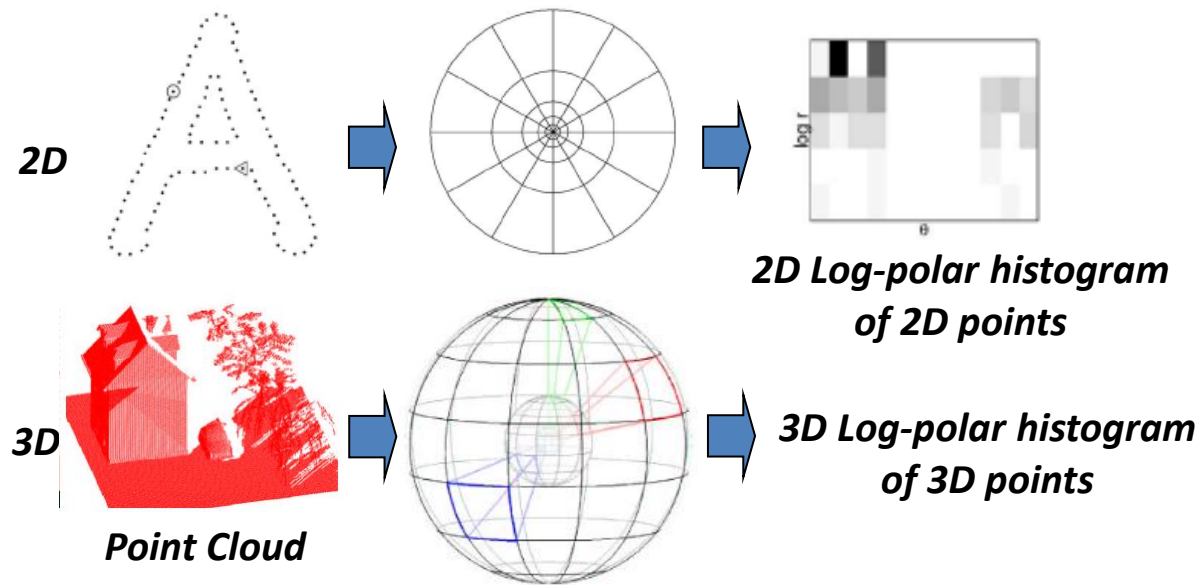
(courtesy of Johnson & Hebert)

- Rotation and translation invariant, not scale invariant
- Appreciates uniform surface sampling
- Variants: compressed-SI (PCA)
- ***pcl::SpinImageEstimation***



Effect of bin size (courtesy of Johnson & Hebert)

- **3DSC** [Frome ECCV04]: extension of the Shape Contexts approach [Belongie et al. PAMI02] to the 3D domain (`pcl::ShapeContext3DEstimation`)
- Each point is accumulated in the 3D bin it falls in, being weighted proportionally to the local point cloud density around the bin and to the bin volume
- No unique local Reference Frame -> L descriptions for each feature (L : number of azimuth bins)




- **Unique Shape Context (USC)** [Tombari 3DOR10]: a unique local RF is plugged in to orient univocally the 3D grid (`pcl::UniqueShapeContext`)
- Hence, only one description is needed for each feature point, decreasing the number of possible mismatches (spurious correspondences) during the matching stage.

- PFH [Rusu08] computes 3 values for each pair in the neighbourhood
 - Complexity $O(k^2)$, extremely slow.
- ***pcl::PFHEstimation***
- For each pair, it computes a LRF u - v - w centred on one point p_s as
 - The normal $u = n_s$
 - The cross product between n_s and the vector $(p_t - p_s)$ $v = n_s \times (p_t - p_s)$
 - The cross product between the previous vectors $w = u \times v$
- Then, it computes and accumulates

$$\alpha = \arccos(v \cdot n_t)$$


$$\phi = \arccos\left(u \cdot \frac{(p_t - p_s)}{\|p_t - p_s\|_2}\right)$$

$$\theta = \arctan(w \cdot n_t, u \cdot n_t)$$






- FPFH [Rusu09]: approximation of PFH with linear complexity in the number of neighbors
 - Compute SPFH (Simplified PFH) between the keypoint and every neighbor
 - Combine the weighted SPFHs to form the final Fast PFH

$$FPFH(p_i) = SPFH(p_i) + \frac{1}{k} \sum_{j=1}^k \frac{1}{\omega_j} SPFH(p_j)$$




- `pcl::FPFHEstimation`, `pcl::FPFHEstimationOMP`

- Signatures of Histograms of Orientations [Tombari10]
- Inspired by SIFT: computation of a geometric coarsely localized local set of histograms of first-order derivatives.
- The local support is partitioned by means of a spherical grid
- For each volume of the grid, an histogram of the cosines of the angle θ_i between the normal at each point and the normal at the feature point is computed.
- Quadrilinear interpolation to smooth out quantization distortions
- Normalization of the descriptor for robustness towards point density variations
- ***pcl::SHOTEstimation, pcl::SHOTEstimationOMP***



- SHOT for RGB-D data [Tombari11] deploys
 - Shape, as the SHOT descriptor
 - Texture, as histograms in the *Lab* space
 - Pairs of *Lab* triplets (center point and its neighbor) can be compared using specific metrics (CIE94, CIE2000, ..), although the L1-norm proved to be a good trade-off
- ***pcl::SHOTColorEstimation*, *pcl::SHOTColorEstimationOMP***



```
pcl::PointCloud<pcl::SHOT352>::Ptr descriptors (new
    pcl::PointCloud<pcl::SHOT352>());

pcl::SHOTEstimationOMP<PointType, NormalType, DescriptorType> describer;





describer.setRadiusSearch (support_radius);
describer.setInputCloud (keypoints);
describer.setInputNormals (normals);
describer.setSearchSurface (cloud);

describer.compute (*descriptors);
```


Summing up..

Method	Category	Unique LRF	Texture
Struct. Indexing [Stein92]	Signature	No	No
PS [Chua97]	Signature	No	No
3DPF [Sun01]	Signature	No	No
3DGSS [Novatnack08]	Signature	No	No
KPQ [Mian10]	Signature	No	No
3D-SURF [Knopp10]	Signature	Yes	No
SI [Johnson99]	Histogram	RA	No
LSP [Chen07]	Histogram	RA	No
3DSC [Frome04]	Histogram	No	No
ISS [Zhong09]	Histogram	No	No
USC [Tombari10]	Histogram	Yes	No
PFH [Rusu08]	Histogram	RA	No
FPFH [Rusu09]	Histogram	RA	No
Tensor [Mian06]	Histogram	No	No
RSD [Marton11]	Histogram	RA	No
HKS [Sun09]	Other	-	No
MeshHoG [Zaharescu09]	Hybrid	Yes	Yes
SHOT [Tombari10]	Hybrid	Yes	Yes

- Taxonomy for global descriptors [Akgul09]
- **Histogram-based:** accumulators of local or global features
 - Robustness, paid off with less descriptiveness
 - Shape Distributions [Osada02], 3D Shape Histograms [Ankerst99], Orientation Histograms [Horn84], **Viewpoint Feature Histogram (VFH) [Rusu10]**, **Clustered-VFH [Aldoma11]**, **OUR-CVFH [Aldoma12]**
- **Transform-based:** Transform geometric information in a domain where representation is compact and invariant
 - Compact descriptors by retaining only a subset of (eg. the first) coefficients
 - 3D Fourier Transform [Dutagaci05], Angular Radial Tr. [Ricard05], 3D Radon Tr. [Daras04], Spherical Harmonics [Kazhdan03], wavelets [Laga06]
- **2D view-based:** 3D surface is transformed into a set of 2D projections (range maps)
 - 2D image descriptors are computed on each 2D view
 - Fourier descriptors [Vranic 04], Zernike moments [Chen03], SIFT [Ohbuchi08], SURF, ..
- **Graph-based:** A graph is built out of the surface
 - Transform the graph into a vector-based numerical description
 - topology-based [Hilaga01], Reeb graph [Tung05], skeleton-based [Sundar03]



- Viewpoint Feature Histogram [Rusu 10]
 - Each 3D model is rendered into different views
 - Each view provides one descriptor
 - Explicitly encodes the viewpoint from where the surface was captured/sensed
 - Based on Point Feature Histogram (PFH)
 - For each point pair **(p_i, p_c)**:
 - Compute a LRF for the **centroid**
 - $u = n_c$
 - $v = \frac{p_i - p_c}{\|p_i - p_c\|} \times u$
 - $w = u \times v$
 - $\alpha = \arccos(v \cdot n)$
 - $\phi = \arccos\left(u \cdot \frac{p_i - p_c}{\|p_i - p_c\|}\right)$
 - $\theta = \text{atan2}(w \cdot n, u \cdot n)$




- Descriptor is built with:


- 3 “PFH” angular values (α , θ , Φ) wrt. centroid (45 bins each)
- 1 shape distribution-like component wrt. centroid (45 bins):

$$SDC = \frac{(p_c - p_i)^2}{\max((p_c - p_i)^2)}$$

- 1 angular value (angle between normal and central view direction – α) (128 bins)







FPFH (α , θ , Φ) Shape dist. Viewpoint (α)
(45*3 bins) (45 bins) (128 bins)





- ***pcl::VFHEstimation***

- The reference frame from VFH is sensitive to missing parts in the surface
- Clustered VFH (CVFH) [Aldoma 11]
 - Perform a further smooth region segmentation on each view
 - Apply a VFH descriptor on each connected component (*cluster*) – no normalization to encode the real size of the object
- VFH, CVFH et al. still present invariance (ambiguity) on the roll angle
- **Camera Roll Histogram** to determine a full 6DO pose
 - distribution of normal angles of all points projected on the camera plane
 - «shift» along roll angle computed by matching CRHs
- *pcl::CVFHEstimation*



- Improvement [Aldoma DAGM13] to CVFH with
 - More robust to missing data
 - A descriptor for each smooth region composing the descriptor)
- One Local-Global Reference Frame for each cluster
 - Locally: compute principal directions
 - Globally: sign disambiguation
- RF splits space in octants:
 - For each octant, Shape distribution (D1) (13 bins)
- Final descriptor:
 - Octant-based shape dist. and color hist.
 - “Global” normal distribution (CVFH) (45x3 els.)
 - Viewpoint (CVFH) (64 els.) (half wrt. VFH/CVFH)
 - Overall size: $13 \times 8 + 199 = 303$
- ***pcl::OURCVFHEstimation*** (currently only trunk)





- Problem: find the kNN of a n-dimensional query vector \mathbf{q} within a set of m candidates (same size)
 - Variant: find all neighbors within an hypersphere of radius r centered on \mathbf{q}
- To speed up the brute force, fast indexing schemes
 - Kd-tree [Freidman77]
 - Hierarchical k-means tree [Fukunaga75]
 - Locality Sensitive Hashing (LSH) [Andoni06]
- Kd-tree slows down at high dimensions (too many nodes, long exploration time), need for approximate kd-tree search
 - Best Bin First [Beis97]
 - Randomized kd-tree [Silpa-Anan08]
 - FLANN [Muja09]
- Example: `pcl::KdTreeFLANN<pcl::SHOT352> matcher;` (in `pcl_kdtree` module)
(also have a look at `pcl::search::FlannSearch`)

- Thanks to: Samuele Salti, Aitor Aldoma